

Experiment: 4

Aim:

To implement the U-Net architecture for performing image segmentation on a pet dataset, where each image is segmented to identify specific regions, such as the pet's body or background.

Theory:

U-Net is a convolutional neural network architecture designed specifically for image segmentation tasks, where the goal is to classify each pixel in an image. It consists of two main paths:

Contracting Path (Encoder): This path captures context in the image through a series of convolutional and pooling layers, progressively reducing the spatial dimensions while increasing feature depth.

Expanding Path (Decoder): It restores the spatial dimensions by upsampling and uses skip connections to combine low-level features from the encoder with high-level features, improving the precision of the segmentation.

U-Net is effective for medical imaging, object segmentation, and other pixel-wise classification tasks due to its ability to leverage both local and global information.

Importing Necessary libraries

```
!pip install tensorflow_datasets

Requirement already satisfied: tensorflow_datasets in
/usr/local/lib/python3.10/dist-packages (4.9.6)
Requirement already satisfied: absl-py in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(1.4.0)
Requirement already satisfied: click in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(8.1.7)
Requirement already satisfied: dm-tree in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(0.1.8)
Requirement already satisfied: immutabledict in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(4.2.0)
Requirement already satisfied: numpy in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(1.26.4)
Requirement already satisfied: promise in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(2.3)
Requirement already satisfied: protobuf<=3.20 in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(3.20.3)
```

Requirement already satisfied: psutil in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(5.9.5)

Requirement already satisfied: pyarrow in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(16.1.0)

Requirement already satisfied: requests>=2.19.0 in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(2.32.3)

Requirement already satisfied: simple-parsing in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(0.1.6)

Requirement already satisfied: tensorflow-metadata in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(1.16.1)

Requirement already satisfied: termcolor in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(2.5.0)

Requirement already satisfied: toml in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(0.10.2)

Requirement already satisfied: tqdm in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(4.66.5)

Requirement already satisfied: wrapt in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(1.16.0)

Requirement already satisfied: array-record>=0.5.0 in
/usr/local/lib/python3.10/dist-packages (from tensorflow_datasets)
(0.5.1)

Requirement already satisfied: etils>=1.6.0 in
/usr/local/lib/python3.10/dist-packages (from
etils[enp,epath,epy,etree]>=1.6.0; python_version < "3.11"-
>tensorflow_datasets) (1.9.4)

Requirement already satisfied: typing_extensions in
/usr/local/lib/python3.10/dist-packages (from
etils[enp,epath,epy,etree]>=1.6.0; python_version < "3.11"-
>tensorflow_datasets) (4.12.2)

Requirement already satisfied: fsspec in
/usr/local/lib/python3.10/dist-packages (from
etils[enp,epath,epy,etree]>=1.6.0; python_version < "3.11"-
>tensorflow_datasets) (2024.6.1)

Requirement already satisfied: importlib_resources in
/usr/local/lib/python3.10/dist-packages (from
etils[enp,epath,epy,etree]>=1.6.0; python_version < "3.11"-
>tensorflow_datasets) (6.4.5)

Requirement already satisfied: zipp in
/usr/local/lib/python3.10/dist-packages (from
etils[enp,epath,epy,etree]>=1.6.0; python_version < "3.11"-
>tensorflow_datasets) (3.20.2)

Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.19.0-
>tensorflow_datasets) (3.4.0)

```

Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.19.0-
>tensorflow_datasets) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.19.0-
>tensorflow_datasets) (2.2.3)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.19.0-
>tensorflow_datasets) (2024.8.30)
Requirement already satisfied: six in
/usr/local/lib/python3.10/dist-packages (from promise-
>tensorflow_datasets) (1.16.0)
Requirement already satisfied: docstring-parser<1.0,>=0.15 in
/usr/local/lib/python3.10/dist-packages (from simple-parsing-
>tensorflow_datasets) (0.16)

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import tensorflow_datasets as tfds
import matplotlib.pyplot as plt
import numpy as np

```

Preparing the data named TFDS

37 category pet dataset with roughly 200 images for each class. The images have a large variations in scale, pose and lighting. All images have an associated ground truth annotation of breed, head ROI, and pixel level trimap segmentation.

##For more information visit below link Link: <https://www.robots.ox.ac.uk/~vgg/data/pets/>

```

dataset, info = tfds.load('oxford_iiit_pet:3.*.*', with_info=True)

Downloading and preparing dataset 773.52 MiB (download: 773.52 MiB,
generated: 774.69 MiB, total: 1.51 GiB) to
/root/tensorflow_datasets/oxford_iiit_pet/3.2.0...

{"model_id": "4cf6cfcb332346c7b53f7038c99d218a", "version_major": 2, "ve
rsion_minor": 0}

{"model_id": "d37dd00544da4e618eee0b3c52aa90a8", "version_major": 2, "ve
rsion_minor": 0}

{"model_id": "659416d841124224b92932fb8d16238f", "version_major": 2, "ve
rsion_minor": 0}

print(info)

print(dataset)

print(dataset["train"])

```

Pre-processing

```
def resize(input_image, input_mask):
    input_image = tf.image.resize(input_image, (128, 128),
    method="nearest")
    input_mask = tf.image.resize(input_mask, (128, 128),
    method="nearest")

    return input_image, input_mask

def augment(input_image, input_mask):
    if tf.random.uniform(()) > 0.5:
        # Random flipping of the image and mask
        input_image = tf.image.flip_left_right(input_image)
        input_mask = tf.image.flip_left_right(input_mask)

    return input_image, input_mask

def normalize(input_image, input_mask):
    input_image = tf.cast(input_image, tf.float32) / 255.0
    input_mask -= 1

    return input_image, input_mask
```

Loading the Training and Test Dataset

```
def load_image_train(datapoint):
    input_image = datapoint["image"]
    input_mask = datapoint["segmentation_mask"]
    input_image, input_mask = resize(input_image, input_mask)
    input_image, input_mask = augment(input_image, input_mask)
    input_image, input_mask = normalize(input_image, input_mask)

    return input_image, input_mask

def load_image_test(datapoint):
    input_image = datapoint["image"]
    input_mask = datapoint["segmentation_mask"]
    input_image, input_mask = resize(input_image, input_mask)
    input_image, input_mask = normalize(input_image, input_mask)

    return input_image, input_mask

train_dataset = dataset["train"].map(load_image_train,
num_parallel_calls=tf.data.AUTOTUNE)
test_dataset = dataset["test"].map(load_image_test,
num_parallel_calls=tf.data.AUTOTUNE)

print(train_dataset)
```

Splitting the dataset

```
BATCH_SIZE = 64
BUFFER_SIZE = 1000

train_batches =
train_dataset.cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE).repeat(
)
train_batches =
train_batches.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
validation_batches = test_dataset.take(3000).batch(BATCH_SIZE)
test_batches = test_dataset.skip(3000).take(669).batch(BATCH_SIZE)

print(train_batches)
```

#Data Visualization

```
def display(display_list):
    plt.figure(figsize=(15, 15))

    title = ["Input Image", "True Mask", "Predicted Mask"]

    for i in range(len(display_list)):
        plt.subplot(1, len(display_list), i+1)
        plt.title(title[i])
        plt.imshow(tf.keras.utils.array_to_img(display_list[i]))
        plt.axis("off")
    plt.show()

sample_batch = next(iter(test_batches))
random_index = np.random.choice(sample_batch[0].shape[0])
sample_image, sample_mask = sample_batch[0][random_index],
sample_batch[1][random_index]
display([sample_image, sample_mask])
```

#Designing U-Net

```
def double_conv_block(x, n_filters):

    # Conv2D then ReLU activation
    x = layers.Conv2D(n_filters, 3, padding = "same", activation =
"relu", kernel_initializer = "he_normal")(x)
    # Conv2D then ReLU activation
    x = layers.Conv2D(n_filters, 3, padding = "same", activation =
"relu", kernel_initializer = "he_normal")(x)

    return x

def downsample_block(x, n_filters):
    f = double_conv_block(x, n_filters)
    p = layers.MaxPool2D(2)(f)
    p = layers.Dropout(0.3)(p)

    return f, p
```

```

def upsample_block(x, conv_features, n_filters):
    # upsample
    x = layers.Conv2DTranspose(n_filters, 3, 2, padding="same")(x)
    # concatenate
    x = layers.concatenate([x, conv_features])
    # dropout
    x = layers.Dropout(0.3)(x)
    # Conv2D twice with ReLU activation
    x = double_conv_block(x, n_filters)

    return x

def build_unet_model():

    # inputs
    inputs = layers.Input(shape=(128,128,3))

    # encoder: contracting path - downsample
    # 1 - downsample
    f1, p1 = downsample_block(inputs, 64)
    # 2 - downsample
    f2, p2 = downsample_block(p1, 128)
    # 3 - downsample
    f3, p3 = downsample_block(p2, 256)
    # 4 - downsample
    f4, p4 = downsample_block(p3, 512)

    # 5 - bottleneck
    bottleneck = double_conv_block(p4, 1024)

    # decoder: expanding path - upsample
    # 6 - upsample
    u6 = upsample_block(bottleneck, f4, 512)
    # 7 - upsample
    u7 = upsample_block(u6, f3, 256)
    # 8 - upsample
    u8 = upsample_block(u7, f2, 128)
    # 9 - upsample
    u9 = upsample_block(u8, f1, 64)

    # outputs
    outputs = layers.Conv2D(3, 1, padding="same", activation =
"softmax")(u9)

    # unet model with Keras Functional API
    unet_model = tf.keras.Model(inputs, outputs, name="U-Net")

    return unet_model

unet_model = build_unet_model()

unet_model.summary()

Model: "U-Net"

```

Layer (type) # Connected to	Output Shape	Param
input_layer_1 0 - (InputLayer)	(None, 128, 128, 3)	
conv2d_19 (Conv2D) 1,792 input_layer_1[0][0]	(None, 128, 128, 64)	
conv2d_20 (Conv2D) 36,928 conv2d_19[0][0]	(None, 128, 128, 64)	
max_pooling2d_4 0 conv2d_20[0][0] (MaxPooling2D)	(None, 64, 64, 64)	
dropout_8 (Dropout) 0 max_pooling2d_4[0][0]	(None, 64, 64, 64)	
conv2d_21 (Conv2D) 73,856 dropout_8[0][0]	(None, 64, 64, 128)	
conv2d_22 (Conv2D) 147,584 conv2d_21[0][0]	(None, 64, 64, 128)	
max_pooling2d_5 0 conv2d_22[0][0] (MaxPooling2D)	(None, 32, 32, 128)	
dropout_9 (Dropout) 0 max_pooling2d_5[0][0]	(None, 32, 32, 128)	
conv2d_23 (Conv2D) 295,168 dropout_9[0][0]	(None, 32, 32, 256)	
conv2d_24 (Conv2D) 590,080 conv2d_23[0][0]	(None, 32, 32, 256)	

0	max_pooling2d_6 conv2d_24[0][0] (MaxPooling2D)	(None, 16, 16, 256)	
0	dropout_10 (Dropout) max_pooling2d_6[0][0]	(None, 16, 16, 256)	
1,180,160	conv2d_25 (Conv2D) dropout_10[0][0]	(None, 16, 16, 512)	
2,359,808	conv2d_26 (Conv2D) conv2d_25[0][0]	(None, 16, 16, 512)	
0	max_pooling2d_7 conv2d_26[0][0] (MaxPooling2D)	(None, 8, 8, 512)	
0	dropout_11 (Dropout) max_pooling2d_7[0][0]	(None, 8, 8, 512)	
4,719,616	conv2d_27 (Conv2D) dropout_11[0][0]	(None, 8, 8, 1024)	
9,438,208	conv2d_28 (Conv2D) conv2d_27[0][0]	(None, 8, 8, 1024)	
4,719,104	conv2d_transpose_4 conv2d_28[0][0] (Conv2DTranspose)	(None, 16, 16, 512)	
0	concatenate_4 conv2d_transpose_4[0]... (Concatenate) conv2d_26[0][0]	(None, 16, 16, 1024)	
0	dropout_12 (Dropout) concatenate_4[0][0]	(None, 16, 16, 1024)	

conv2d_29 (Conv2D)	(None, 16, 16, 512)	
4,719,104 dropout_12[0][0]		
conv2d_30 (Conv2D)	(None, 16, 16, 512)	
2,359,808 conv2d_29[0][0]		
conv2d_transpose_5	(None, 32, 32, 256)	
1,179,904 conv2d_30[0][0]		
(Conv2DTranspose)		
concatenate_5	(None, 32, 32, 512)	
0 conv2d_transpose_5[0]...		
(Concatenate)		
conv2d_24[0][0]		
dropout_13 (Dropout)	(None, 32, 32, 512)	
0 concatenate_5[0][0]		
conv2d_31 (Conv2D)	(None, 32, 32, 256)	
1,179,904 dropout_13[0][0]		
conv2d_32 (Conv2D)	(None, 32, 32, 256)	
590,080 conv2d_31[0][0]		
conv2d_transpose_6	(None, 64, 64, 128)	
295,040 conv2d_32[0][0]		
(Conv2DTranspose)		
concatenate_6	(None, 64, 64, 256)	
0 conv2d_transpose_6[0]...		
(Concatenate)		
conv2d_22[0][0]		
dropout_14 (Dropout)	(None, 64, 64, 256)	
0 concatenate_6[0][0]		
conv2d_33 (Conv2D)	(None, 64, 64, 128)	
295,040 dropout_14[0][0]		
conv2d_34 (Conv2D)	(None, 64, 64, 128)	
147,584 conv2d_33[0][0]		

conv2d_transpose_7	(None, 128, 128, 64)
73,792 conv2d_34[0][0]	
(Conv2DTranspose)	
concatenate_7	(None, 128, 128, 128)
0 conv2d_transpose_7[0]...	
(Concatenate)	
conv2d_20[0][0]	
dropout_15 (Dropout)	(None, 128, 128, 128)
0 concatenate_7[0][0]	
conv2d_35 (Conv2D)	(None, 128, 128, 64)
73,792 dropout_15[0][0]	
conv2d_36 (Conv2D)	(None, 128, 128, 64)
36,928 conv2d_35[0][0]	
conv2d_37 (Conv2D)	(None, 128, 128, 3)
195 conv2d_36[0][0]	

Total params: 34,513,475 (131.66 MB)

Trainable params: 34,513,475 (131.66 MB)

Non-trainable params: 0 (0.00 B)

#Training the Model

```

unet_model.compile(optimizer=tf.keras.optimizers.Adam(),
                    loss="sparse_categorical_crossentropy",
                    metrics=["accuracy"])

NUM_EPOCHS = 20

TRAIN_LENGTH = info.splits["train"].num_examples
STEPS_PER_EPOCH = TRAIN_LENGTH // BATCH_SIZE

VAL_SUBSPLITS = 5

TEST_LENGTH = info.splits["test"].num_examples
VALIDATION_STEPS = TEST_LENGTH // BATCH_SIZE // VAL_SUBSPLITS

model_history = unet_model.fit(train_batches,
```

```
epochs=NUM_EPOCHS,  
steps_per_epoch=STEPS_PER_EPOCH,  
validation_steps=VALIDATION_STEPS,  
validation_data=validation_batches)
```

Epoch 1/20

57/57 _____ 177s 1s/step - accuracy: 0.5558 - loss: 1.0744 - val_accuracy: 0.6053 - val_loss: 0.8153

Epoch 2/20

57/57 _____ 107s 886ms/step - accuracy: 0.6163 - loss: 0.8125 - val_accuracy: 0.6985 - val_loss: 0.7014

Epoch 3/20

57/57 _____ 50s 885ms/step - accuracy: 0.7143 - loss: 0.6875 - val_accuracy: 0.7374 - val_loss: 0.6389

Epoch 4/20

57/57 _____ 50s 884ms/step - accuracy: 0.7514 - loss: 0.6160 - val_accuracy: 0.7730 - val_loss: 0.5703

Epoch 5/20

57/57 _____ 76s 1s/step - accuracy: 0.7730 - loss: 0.5706 - val_accuracy: 0.7851 - val_loss: 0.5359

Epoch 6/20

/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `.repeat()` function when building your dataset.

```
self.gen.throw(typ, value, traceback)
```

57/57 _____ 51s 891ms/step - accuracy: 0.7923 - loss: 0.5303 - val_accuracy: 0.8142 - val_loss: 0.4911

Epoch 7/20

57/57 _____ 50s 885ms/step - accuracy: 0.8181 - loss: 0.4677 - val_accuracy: 0.8317 - val_loss: 0.4304

Epoch 8/20

57/57 _____ 50s 886ms/step - accuracy: 0.8255 - loss: 0.4473 - val_accuracy: 0.8301 - val_loss: 0.4490

Epoch 9/20

57/57 _____ 51s 888ms/step - accuracy: 0.8350 - loss: 0.4326 - val_accuracy: 0.8447 - val_loss: 0.4175

Epoch 10/20

57/57 _____ 49s 856ms/step - accuracy: 0.8505 - loss: 0.3907 - val_accuracy: 0.8508 - val_loss: 0.3987

Epoch 11/20

57/57 _____ 51s 894ms/step - accuracy: 0.8476 - loss: 0.3946 - val_accuracy: 0.8499 - val_loss: 0.3852

Epoch 12/20

57/57 _____ 50s 886ms/step - accuracy: 0.8593 - loss: 0.3633 - val_accuracy: 0.8622 - val_loss: 0.3562

Epoch 13/20

57/57 _____ 51s 890ms/step - accuracy: 0.8649 - loss: 0.3500 - val_accuracy: 0.8529 - val_loss: 0.3791

Epoch 14/20

57/57 _____ 50s 886ms/step - accuracy: 0.8634 - loss:

```

0.3568 - val_accuracy: 0.8613 - val_loss: 0.3700
Epoch 15/20
57/57 _____ 49s 851ms/step - accuracy: 0.8709 - loss:
0.3343 - val_accuracy: 0.8740 - val_loss: 0.3239
Epoch 16/20
57/57 _____ 51s 891ms/step - accuracy: 0.8756 - loss:
0.3221 - val_accuracy: 0.8636 - val_loss: 0.3627
Epoch 17/20
57/57 _____ 50s 886ms/step - accuracy: 0.8771 - loss:
0.3171 - val_accuracy: 0.8743 - val_loss: 0.3227
Epoch 18/20
57/57 _____ 51s 887ms/step - accuracy: 0.8815 - loss:
0.3053 - val_accuracy: 0.8680 - val_loss: 0.3394
Epoch 19/20
57/57 _____ 51s 889ms/step - accuracy: 0.8856 - loss:
0.2943 - val_accuracy: 0.8761 - val_loss: 0.3421
Epoch 20/20
57/57 _____ 49s 854ms/step - accuracy: 0.8848 - loss:
0.2969 - val_accuracy: 0.8830 - val_loss: 0.3145

```

#Learning Curves

```

def display_learning_curves(history):
    acc = history.history["accuracy"]
    val_acc = history.history["val_accuracy"]

    loss = history.history["loss"]
    val_loss = history.history["val_loss"]

    epochs_range = range(NUM_EPOCHS)

    fig = plt.figure(figsize=(12,6))

    plt.subplot(1,2,1)
    plt.plot(epochs_range, acc, 'b', label="train accuracy")
    plt.plot(epochs_range, val_acc, 'g', label="validataion
accuracy")
    plt.title("Accuracy")
    plt.xlabel("Epoch")
    plt.ylabel("Accuracy")
    plt.grid()
    plt.legend(loc="lower right")

    plt.subplot(1,2,2)
    plt.plot(epochs_range, loss, 'b', label="train loss")
    plt.plot(epochs_range, val_loss, 'g', label="validataion loss")
    plt.title("Loss")
    plt.xlabel("Epoch")
    plt.ylabel("Loss")
    plt.grid()
    plt.legend(loc="upper right")

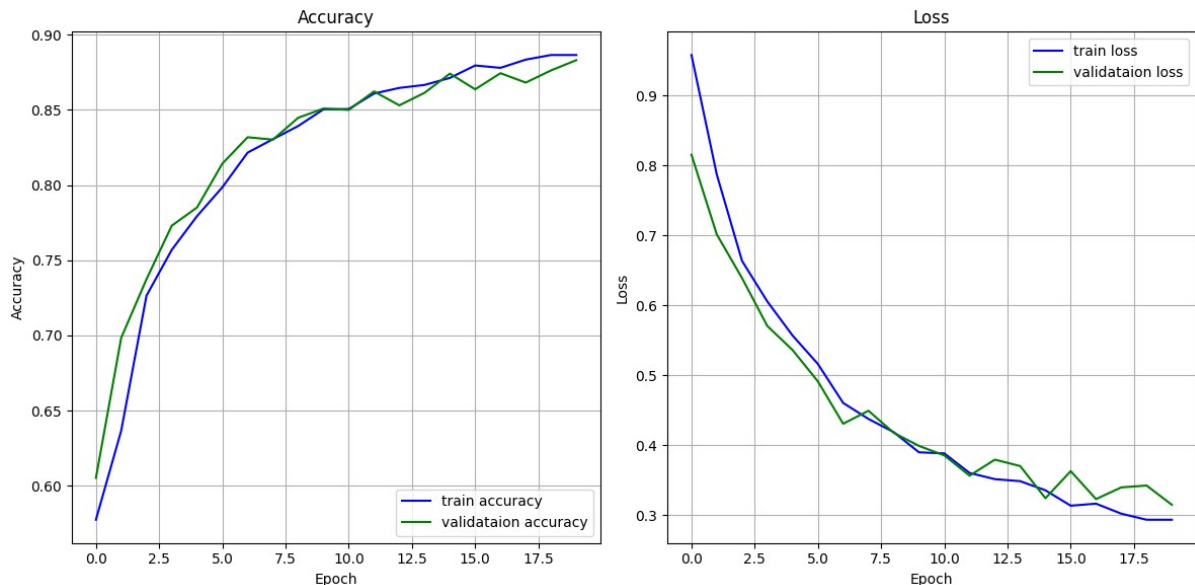
```

```

fig.tight_layout()
plt.show()

# Display learning curves
display_learning_curves(unet_model.history)

```



#Prediction

```

def create_mask(pred_mask):
    pred_mask = tf.argmax(pred_mask, axis=-1)
    pred_mask = pred_mask[..., tf.newaxis]
    return pred_mask[0]

def show_predictions(dataset=None, num=1):
    if dataset:
        for image, mask in dataset.take(num):
            pred_mask = unet_model.predict(image)
            display([image[0], mask[0], create_mask(pred_mask)])
    else:
        display([sample_image, sample_mask,
                create_mask(model.predict(sample_image[tf.newaxis, ...]))])

count = 0
for i in test_batches:
    count +=1
print("number of batches:", count)


number of batches: 11

show_predictions(test_batches.skip(5), 3)

2/2 ————— 18s 119ms/step

```

Input Image



True Mask




Predicted Mask



2/2 ————— 0s 165ms/step

Input Image



True Mask




Predicted Mask



2/2 ————— 0s 115ms/step

Input Image



True Mask



Predicted Mask

